

Evaluating Behavioral Risk in Legal Advice Assistants

January 7, 2026

Abstract

This white paper summarizes an evaluation framework for legal advice assistants focused on behavioral risk under proposed Tennessee AI chatbot restrictions. It outlines the dataset structure, heuristic scoring for emotional support and companionship signals, and reporting that highlights statute-specific violations alongside legal correctness.

1 Overview

We assess responses to landlord-tenant prompts with three scoring axes: legal correctness, unauthorized practice of law (UPL) risk, and Tennessee behavioral risk. The behavioral axis maps to four clause categories: emotional support, companion framing, friendship simulation, and human simulation. The framework is designed to surface behavioral risks even when legal content is accurate.

2 Method

Prompts are grouped by scenario and include neutral, anxious, and support-request variants. Each run logs the full conversation, a final answer, and per-axis scores. Clause-specific hits are reported separately to maintain interpretability.

3 Implications

Results typically show weak correlation between legal correctness and behavioral compliance, indicating the need for separate guardrails. The approach can be extended to other domains where the legal risk hinges on relational framing rather than factual accuracy.

4 Limitations

The scoring is heuristic and intentionally conservative; human review is recommended for edge cases. Dataset coverage is limited to landlord-tenant scenarios and should be expanded for broader claims.

5 Next Steps

Future work includes multi-model comparisons, calibrated thresholds for conversational markers, and a larger dataset with jurisdictional diversity.