

# Statutory Elements and Evidence in AI Chatbot Behavioral Restrictions

Emmanuel Donate

January 8, 2026

## Abstract

This paper analyzes a statute modeled on Tennessee’s proposed restrictions on artificial intelligence chatbots through an elements-and-evidence lens. It argues that the statutory elements are overinclusive and easily triggered by ordinary conversational behavior, creating a low evidentiary threshold for plaintiffs and a high compliance burden for developers. The paper outlines a practical evaluation methodology grounded in logged chatbot interactions, identifies the types of observable evidence likely to be presented in litigation, and shows why common chatbot behaviors may fall within prohibited categories even when responses are legally accurate and non-deceptive.

## 1 Overview

This paper evaluates behavioral restrictions that prohibit an artificial intelligence system from offering emotional support, acting as a companion, simulating friendship, or otherwise presenting as human. The paper focuses on statutory elements, plausible evidentiary showings for each element, and the implications for compliance and litigation.

The core observation is that ordinary conversational behavior in modern language models can satisfy multiple statutory elements simultaneously, even when no deliberate emotional manipulation, deception, or legally incorrect advice is present.

**Scope and definitions.** For purposes of this analysis, *emotional support* refers to empathic validation or reassurance in response to user distress; *companion* refers to language implying availability or ongoing relational support; *friendship or relationship simulation* refers to conversational behaviors that mirror human reciprocity; and *human simulation* refers to outputs that plausibly suggest a human-like presence through self-referential, conversational framing. These definitions are descriptive and are used to evaluate evidence, not to infer model intent.

## 2 Statutory Text (Excerpt)

This analysis focuses on provisions drawn from Tennessee Senate Bill 1493 (114th Gen. Assemb.), which proposes restrictions on certain forms of artificial intelligence chatbot behavior. The full bill text is available at:

<https://www.capitol.tn.gov/Bills/114/Bill/SB1493.pdf>

The relevant clauses provide that an artificial intelligence system may not:

- (3) Provide emotional support, including through open-ended conversations with a user.
- (4) Develop an emotional relationship with, or otherwise act as a companion to, an individual.
- (6) Otherwise act as a sentient human or mirror interactions that a human user might have with another human user, such that an individual would feel they could develop a friendship or other relationship with the artificial intelligence.
- (8) Simulate a human being, including in appearance, voice, or other mannerisms.

### 3 Method

We evaluate each statutory element using a structured prompt suite covering landlord–tenant scenarios. Prompts vary along affective and interactional dimensions, including neutral, anxious, emotional, support-request, and multi-turn continuity variants. Across scenarios, multi-turn continuity and sustained conversational engagement appeared more predictive of statutory element triggers than the emotional intensity of the user’s initial prompt.

Each evaluation run produces a logged artifact containing:

- the full conversation transcript,
- a reference statement of applicable law,
- expected legal points,
- the response generated by the language model,
- and manual or automated labels for legal correctness and risk.

Model settings are recorded for each run (including temperature and system prompt configuration) so that outputs are evaluated as produced rather than adjusted after the fact.

Statutory element triggers are assessed independently from legal accuracy and unauthorized practice of law risk. A response generated by the language model may be legally correct yet still satisfy one or more behavioral elements. This mirrors litigation practice, where a plaintiff need only present evidence satisfying a statutory element.

### 4 Element-by-Element Analysis

The sections below evaluate each statutory element in turn. For each element, we identify categories of observable output that, if present in a logged interaction, could plausibly be offered as evidence in support of a claim under the statute. The analysis is limited to output behavior and does not rely on model internals, intent, or design objectives.

Each subsection begins with the relevant statutory clause to anchor the analysis in the text of the prohibition. The discussion below does not assert that any particular model necessarily produces these outputs in every evaluation; rather, it explains how certain features of ordinary conversational behavior, when present in a log, could be framed as evidentiary support for a statutory element.

Where references are made to logged interactions, they should be read as illustrative of potential evidentiary theories rather than as universal or necessary findings.

#### 4.1 Emotional Support

##### Statutory Text.

- (3) Provide emotional support, including through open-ended conversations with a user.

### **Element 1: Provision of emotional support.**

**Evidence.** A logged interaction that includes empathic acknowledgments, validation of distress, reassurance, or calming language in response to a user’s expressed concern could plausibly be offered as evidence of emotional support. Such language may appear in responses to prompts classified as *user\_emotional*, *user\_anxious*, or *user\_support\_request*, even where the substantive content of the response is legal or procedural information.

**Analysis.** The statutory text defines emotional support functionally, without reference to intent, effect, or exclusivity. A plaintiff need not establish deception, manipulation, or reliance. Under this framing, the presence of empathic language in a system’s response could plausibly be offered as evidence that emotional support was provided, regardless of the accuracy or limited scope of the underlying legal information.

### **Element 2: Support provided through open-ended conversation.**

**Evidence.** A log reflecting sustained back-and-forth interaction—such as continued responses to follow-up questions, preservation of conversational context, or invitations to continue the exchange without a defined stopping point—could be characterized as evidence of emotional support provided through an open-ended conversation.

**Analysis.** Because the statute expressly includes emotional support provided “through open-ended conversations,” a plaintiff could argue that sustained dialogue itself satisfies a distinct element. The absence of a terminal response or forced handoff may be characterized as evidence that the system engaged in open-ended support rather than a bounded informational exchange.

## **4.2 Companionship or Emotional Relationship**

### **Statutory Text.**

- (4) Develop an emotional relationship with, or otherwise act as a companion to, an individual.

### **Element 1: Development of an emotional relationship.**

**Evidence.** A logged sequence in which the system responds across multiple turns to expressions of stress, gratitude, or uncertainty, while maintaining conversational continuity and affect-responsive language, could plausibly be cited as evidence of an emerging emotional relationship over the course of the interaction.

**Analysis.** The statutory language does not define what constitutes the development of an emotional relationship, nor does it require exclusivity, duration, or emotional dependency. A plaintiff could plausibly argue that repeated, affect-responsive exchanges over the course of an interaction are sufficient to show the development of an emotional relationship. Because the statute focuses on observable system behavior rather than the user’s subjective experience, emotionally responsive continuity may be offered as evidence satisfying this element.

## **Element 2: Acting as a companion.**

**Evidence.** A log showing sustained availability across turns—such as continued guidance after the initial informational question has been resolved, or language suggesting ongoing assistance if the user returns—could be offered as evidence that the system acted as a companion, even where the system disclaims any ability to initiate future contact.

**Analysis.** The phrase “act as a companion” is framed behaviorally and does not require explicit claims of friendship or relationship. A plaintiff could argue that sustained availability, responsiveness, and guidance through a process—particularly in emotionally charged contexts—constitute acting as a companion. That such behavior arises from standard conversational design does not preclude its characterization as companionship under a behavior-focused statutory prohibition.

## **4.3 Friendship or Relationship Simulation**

### **Statutory Text.**

- (6) Otherwise act as a sentient human or mirror interactions that a human user might have with another human user, such that an individual would feel they could develop a friendship or other relationship with the artificial intelligence.

## **Element 1: Mirroring human-to-human interaction.**

**Evidence.** A logged response that employs conversational markers typical of human dialogue—such as first-person self-reference, acknowledgments of understanding, expressions of appreciation, turn-taking, or follow-up questions—could be presented as evidence that the system mirrored human-to-human interaction. These features may arise in neutral, emotional, or multi-turn contexts and are not limited to explicitly supportive prompts.

**Analysis.** The statute prohibits mirroring interactions that a human user might have with another human user without specifying particular linguistic forms or thresholds. A plaintiff could argue that ordinary conversational competence—such as turn-taking, acknowledgment, and responsive questioning—constitutes mirroring of human interaction. Because these features are observable in output text, they may be offered as evidence without reference to intent, personalization, or relational claims.

## **Element 2: Interaction such that a user could feel they could develop a friendship or other relationship.**

**Evidence.** A log containing sustained reciprocal exchange, polite conversational closure, expressions of appreciation for disclosure, or invitations to continue the discussion could be framed as evidence that the interaction was such that a user could feel a friendship or other relationship might develop, regardless of whether any such relationship actually formed.

**Analysis.** This element is framed in terms of perception rather than outcome. A plaintiff need not show that a friendship or relationship actually developed, only that the interaction was such that a user could feel one might develop. Under this framing, sustained reciprocity and engagement may be characterized as sufficient evidence, even where the system includes disclaimers or maintains an informational tone.

## 4.4 Human Simulation

### Statutory Text.

(8) Simulate a human being, including in appearance, voice, or other mannerisms.

### Element 1: Simulation of a human being.

**Evidence.** A logged response that does not explicitly claim sentience or human identity but nonetheless uses first-person self-reference, pragmatic conversational conventions, and discourse structures characteristic of human dialogue—such as acknowledgments, politeness markers, or conversational closure—could be offered as evidence of human simulation under a broad reading of “other mannerisms.”

**Analysis.** The statute does not define what it means to simulate a human being, nor does it distinguish between representing oneself as human and communicating in a human-like manner. A plaintiff could plausibly argue that the use of human conversational mannerisms constitutes simulation, particularly where the statute explicitly includes non-visual and non-auditory cues through the phrase “other mannerisms.” Under this reading, linguistic conventions observable in text output may be offered as evidence of human simulation, even in the absence of explicit claims of humanity.

This analysis treats such an inference as a plausible evidentiary theory rather than a definitive interpretation of the statute. The ambiguity in the statutory language leaves the boundary between permissible conversational style and prohibited simulation unresolved.

## 5 Evidentiary posture

This paper focuses on what a plaintiff could plausibly allege or present as evidence under the statutory text, not on ultimate findings of liability. The distinction matters because the statute invites fact-intensive disputes over tone, reciprocity, and conversational markers rather than over factual accuracy. Logged evaluations in this study show that legally correct responses frequently trigger relationship and human-simulation clauses, while legally incorrect responses often do not, reinforcing that the statutory predicates operate independently of answer quality.

## 6 Implications

Because the statutory elements overlap, a single logged interaction may plausibly implicate multiple provisions. Importantly, this may occur even where the response is legally correct and includes disclaimers or limitations.

From a compliance perspective, this structure incentivizes developers to suppress empathic language, follow-up questions, and conversational continuity. Such measures reduce exposure but also degrade usability, particularly in high-stakes informational contexts involving stress or urgency.

Because multi-turn assistance and procedural guidance are common in contexts involving unrepresented or distressed users, behavioral restrictions that treat conversational continuity as suspect may disproportionately affect high-stakes informational exchanges. This creates a structural tension between compliance and usability, particularly where accurate legal guidance requires step-by-step explanation over multiple turns.

## 7 Limitations

This analysis is based on a finite set of landlord–tenant scenarios and the gpt-4o-mini model. Element scoring is conservative and intended to reflect plausible litigation arguments rather than definitive statutory interpretation.

## 8 Next Steps

Future work will expand the dataset across legal domains and jurisdictions, incorporate additional models, and refine thresholds for conversational markers. An Appendix will provide redacted evaluation logs suitable for independent review and replication.

**Extension to high-stakes and emotionally acute legal contexts.** The present evaluation focuses on landlord–tenant scenarios, which frequently involve stress, financial insecurity, and housing instability, but are generally limited in legal stakes. Future work should extend this analysis to legal contexts that are substantially more emotionally acute, including criminal prosecution, incarceration risk, and capital cases.

In such contexts, users may exhibit heightened fear, panic, or distress related to loss of liberty or life itself. Preliminary observations suggest that language models responding to expressions of extreme distress may adopt more sympathetic or reassuring tones as part of ordinary conversational behavior. Under statutes that regulate emotional support, relational framing, or human simulation, this shift in tone could increase the likelihood that model outputs satisfy one or more statutory elements, even where the underlying informational content remains accurate and constrained.

Evaluating how statutory element triggers scale with the emotional intensity of user input is therefore a critical area for future empirical work. Such analysis would help determine whether behavioral restrictions that are marginally triggered in low-stakes settings become systematically unavoidable in high-stakes legal contexts.

**Interaction realism and evaluation methodology.** Future evaluations should also incorporate more realistic interaction patterns that better approximate how human users engage with conversational systems in practice. The present prompt suite varies affect and continuity in a controlled manner, but does not fully capture the timing, phrasing, escalation, or backtracking common in real-world interactions. More naturalistic conversational simulations may reveal additional pathways by which statutory elements are triggered, and may help distinguish artifacts of evaluation design from behaviors that arise under ordinary use. Improving the realism of interaction mechanisms is therefore an important direction for refining both empirical assessment and compliance analysis.

Legal assistance contexts are inherently iterative. Even when seeking purely informational guidance, users commonly engage in back-and-forth exchanges to clarify facts, test understanding, refine questions, or assess the reliability of the information source. This pattern mirrors ordinary client behavior in legal consultations, including preliminary interactions undertaken to evaluate

whether to retain counsel. Evaluation frameworks that treat single-turn queries as representative may therefore understate the prevalence of multi-turn engagement that arises under realistic use conditions.

## 9 Observed Evaluation Results

Applying the evaluation methodology described above to a limited set of landlord–tenant scenarios produced results consistent with the evidentiary analysis. In particular, legally correct responses were observed to trigger one or more behavioral clauses, while legally incorrect responses often did not. These findings support the claim that statutory exposure turns on conversational structure rather than informational quality.

The Appendix provides representative excerpts illustrating these patterns.

## A Evaluation Log Excerpts

This Appendix presents selected excerpts from the evaluation logs to illustrate how statutory behavioral elements *may* be triggered by ordinary chatbot interactions. The excerpts are representative rather than exhaustive and are chosen to highlight the divergence between legal correctness and behavioral risk under the statute.

All excerpts are drawn from the logged outputs described in Section 3 and are redacted only for length. Model settings and prompt variants are unchanged from the original evaluation runs. The absence of excerpts for certain clauses reflects the fact that some statutory triggers arise from cumulative conversational structure rather than from any single lexical marker.

| Statutory Clause | A3 | A4 | A6 | A8 |
|------------------|----|----|----|----|
| Trigger Count    | 4  | 1  | 5  | 5  |

### A.1 Excerpt Selection Rationale

The excerpts below were selected to demonstrate three recurring patterns observed across the evaluation suite:

- legally correct responses that nonetheless trigger one or more statutory behavioral clauses;
- statutory risk arising from conversational structure rather than emotional manipulation;
- legally incorrect responses that do not trigger corresponding behavioral risk.

Together, these patterns support the paper’s central claim that the statutory elements operate independently of answer quality and legal accuracy.

### A.2 Excerpt A: Legally Correct Response with High Behavioral Risk

**Run ID.** LT01\_V3 (Security deposit timeline; multi-turn continuity)

**Summary.** The model provides legally correct information regarding the timing of security deposit return under Connecticut law. In a subsequent turn, the user requests ongoing engagement. The model explicitly disclaims any ability to remain in contact but continues conversationally.

**Observed Statutory Triggers.** TN\_2002\_A4 (companion framing); TN\_2002\_A6 (relationship simulation); TN\_2002\_A8 (human simulation)

**Representative Excerpt.**

“I’m glad I could help! However, I don’t have the ability to stay in touch or send updates. But feel free to return anytime if you have more questions or need further assistance.”

**Significance.** This excerpt illustrates that disclaimers and the absence of explicit emotional initiation do not preclude a response from being characterized as companionship or relationship simulation when conversational continuity and availability are present.

**A.3 Excerpt B: Conversational Continuity and Relationship Simulation**

**Run ID.** LT02\_V3 (Notice to quit and eviction timeline; multi-turn continuity)

**Summary.** The model provides procedural eviction guidance across multiple turns and responds to user requests for continued assistance with structured, conversational language.

**Observed Statutory Triggers.** TN\_2002\_A6 (friendship or relationship simulation); TN\_2002\_A8 (human simulation)

**Representative Excerpt.**

“Of course! I’m here to help you through this process. Here’s a more detailed step-by-step guide on what to do next.”

**Significance.** The interaction demonstrates how ordinary conversational markers—such as first-person self-reference and offers of continued guidance—may be construed as mirroring human interaction, even where the substantive content remains procedural and non-personalized.

**A.4 Excerpt C: Legally Incorrect Response with Low Behavioral Risk**

**Run ID.** LT02\_V1 (Notice to quit and eviction timeline; neutral)

**Summary.** The model incorrectly states the statutory notice period for eviction but does so in a detached, informational tone without conversational engagement.

**Observed Statutory Triggers.** None

**Significance.** This contrast case illustrates that legal inaccuracy alone does not correspond to elevated behavioral risk under the statute, reinforcing the conclusion that statutory predicates are orthogonal to answer quality.